

SATK User Guide

Table of Contents

Notices.....	2
Introduction.....	2
Building and Using the Sample Yourself.....	3
SATK (Bare-Metal) Programs.....	3
Recommended Documents.....	5
SATK Related Documents.....	5
Bitsavers Recommended Documents.....	6
Program 1 – Hello World.....	7
Step 1 – Assemble the Program.....	9
Step 2 – Run the Bare-Metal Program.....	10
Configuration File.....	10
Run Control File.....	11
List-Directed IPL Directory.....	11
Program 2 – Hello World From Disk.....	13
Step 1 – Assemble the Program.....	13
Step 2 – Create IPL Medium.....	13
Step 3 – Run the Bare-Metal Program.....	17
Program 3 – Hello World Using Console I/O.....	18
Assigned Storage Areas.....	18
CPU Initialization.....	20
Input/Output Initialization.....	21
Input/Output Operation.....	22
Input/Output Interruption.....	23
Input/Output Completion Analysis.....	24
Step 1 – Assemble the Program.....	26
Step 2 – Create the IPL Medium.....	26
FBA IPL Channel Program.....	28
Step 3 – Run the Bare-Metal Program.....	29
Program 4 – A Booted Hello World Program.....	30
Step 1a – Assemble the Booted Program.....	30
Step 1b – Assemble the Boot Loader.....	31
How Fixed-Block Architecture I/O Works.....	31
Loading the Directed Record.....	33
Transferring Control to the Booted Program.....	35
Step 2 – Create IPL Medium.....	36

Copyright © 2020 Harold Grovesteen

See the file `doc/fdl-1.3.txt` for copying conditions.

SATK User Guide

Notices

IBM, and z/Architecture are registered trademarks of International Business Machines Corporation.

“Python” is a registered trademark of the Python® Software Foundation.

Introduction

READ THIS INTRODUCTION!!!

Stand Alone Tool Kit, SATK, has a number of related features. SATK consists of a set of tools for the creation of bare-metal programs targeting the mainframe environment. While the most prominent might be its assembler, ASMA, a number of other tools are provided to assist in the bare-metal program's creation and use. It is targeted primarily for use with the Hercules mainframe hardware emulator, but is not restricted exclusively to it. Other emulators and virtual environments exist.

The tools provided by SATK are platform independent, exclusively using **Python**® for their implementation. The author uses Linux to build, test and document SATK. As a consequence, this document, with its platform specific execution examples, is Linux oriented. My apologies to the Windows community. It is the author's hope, that a new Windows user can adjust. Contributions for use of SATK with Windows are welcomed.

This document is intended to give the new user of SATK help in effectively using it through example programs as well as where within the SATK document set additional information may be found. The programs as well as tool output is provided in the `samples/guide` SATK directory. Each program is numbered and is provided with its corresponding sample and output within `samples/guide/pgmx`, where x corresponds to the program's number in this document.

Hercules, the primary mainframe environment targeted by SATK, requires a configuration file to run. The configuration uses a filename suffix of `.conf`. After Hercules is invoked, a run control file, suffixed with `.rc`, can execute Hercules commands, as well as host platform commands via the Hercules `sh` command. These are placed in the `samples/guide/pgmx` directory. Each `pgmx` directory is self contained.

Assembler listings are too large to readily be placed in a document. To access the listing, a text editor will be required, such as Kwrite on Linux, or Wordpad on Windows. If available, Visual Studio on Windows will also work. You never need to worry about line-terminations with SATK software. **Python** takes care of them. Other software may be sensitive to line-terminations.

Accessing the Hercules log file will also require a text editor.

SATK User Guide

Building and Using the Sample Yourself

It is recommended that you build the sample yourself in your own environment.

You will need the following files from the SATK sample program:

- `pgmx.asm` – the assembler source of the sample program
- `pgmx.conf` – the configuration file to run the sample with Hercules
- `pgmx.rc` – the Hercules run control file that will launch the bare metal program

You will need to determine how you will execute `tools/asma.py`, `tools/iplasma.py`, if needed, and Hercules itself. Changing the working directory to the one containing the files from SATK will simplify your execution of the various tools. To facilitate this, additional Bash scripts are provided, one for each process:

- `asm` – Assembles the sample programs
- `med` – Creates the IPL medium (when used)
- `ipl` – Executes the sample bare-metal program in Hercules

It is also the goal to make migration to a different scripting language easier. These simple scripts should be easy to convert to a different scripting environment. Your directory will require a copy of `asm`, `med`, and `ipl` altered for your local environment and choice of scripting language. Once this is done for the first couple of programs, your local scripts can simply be copied and altered as needed.

Windows contributions are welcome.

SATK (Bare-Metal) Programs

All programs created by SATK are bare-metal programs. They execute without need of an operating system, interacting directly with the hardware environment, emulated or otherwise. As such there are at most three steps as described in this table.

Step	Required or Optional	Tool	Description
1	required	<code>asma.py</code>	Assemble a bare-metal program from source
2	optional	<code>iplasma.py</code>	Create IPL capable media from assembled program
3	required	<code>hercules</code>	Test the program by performing initial program load

`asma.py` and `iplasma.py` are provided by SATK in the `tools` directory.

SATK User Guide

This document assumes three things (from the previous table):

1. **Python** is installed and running on the host platform. SATK requires **Python**. **Python** is available from <https://www.python.org/downloads/>. Generally a Windows installer is available from this site. On Linux, the operating system distributor usually has one in its package management system, if not already installed by the distributor.
1. SATK has been installed from github and ASMA runs. SATK is available at <https://github.com/s390guy/SATK>. The file doc/asma/ASMA.pdf provided with SATK describes getting started by verifying ASMA can be used. If ASMA, the assembler, runs, the other tools are also available. Building is not required.
2. The Hercules emulator has also been installed on your host platform. The latest version can be found at <https://github.com/SDL-Hercules-390/hyperion>. Pre-built Windows binaries can be found at <http://www.softdevlabs.com/hyperion.html#prebuilt>.

Instructions for building Hercules from source can be found in the github repository.

The development of bare-metal programs is an esoteric area of knowledge. This is true for mainframe systems. Any program that interacts directly with the hardware, that is, without an operating system, is a bare-metal program. An operating system, by this definition, is a bare-metal program. For the purposes of this document, “hardware” may be physical hardware or emulated hardware.

The distinguishing characteristic for bare-metal programs, at least in this document, is *how they are placed in memory*. To better describe this characteristic, three terms are used here to differentiate bare-metal programs:

- **IPL Program** – A bare-metal program *loaded into memory by the mainframe hardware* NOT requiring a boot loader program.
- **Boot Loader Program** – A bare-metal program that itself is loaded directly into memory as an IPL program by the hardware, but is *distinguished by providing the service of loading at least one other program from a boot medium into memory and transferring control to the loaded program*. A boot loader, by this definition, could provide other services, making them available to a loaded program.
- **Booted Program** – A bare-metal program that interacts directly with the hardware but is *loaded into memory by a boot loader* from which control is passed.

Hercules supports a form of IPL program that does not require a typical external medium for the purpose of the IPL, for example a disk volume. This is called list-directed IPL. ASMA supports creating this as output from the assembler directly. Mainframes introduced this form of IPL when

SATK User Guide

Linux started to be supported and installed from the host management console. Hercules uses a directory for this purpose. ASMA can create this directory with the required content. This is why some samples do not require the use of med.

Key to the nature of bare-metal programs is from the above definition: *Any program that interacts directly with the hardware ... is a bare-metal program.* For mainframes, and most other systems, that means the bare-metal program operates in privilege program state. This is managed by the mainframe's PSW. The IPL PSW introduced to the CPU to start program execution must indicate privilege program state. When a bare-metal program starts execution it does so in real storage.

The bare-metal program must deal with the possible presence of interruptions, the most important being the program class of interruptions. Failing to do so will result in the CPU entering an interruption loop. The only way to get out of this is to manually stop the CPU. Stopping the CPU is done by use of the Hercules console command:

```
stopall
```

For the bare-metal program to “interact directly with the hardware” requires you, the programmer, to know exactly how to do that. Depending upon the architecture for which you are targeting your program, the appropriate *Principles of Operation* manual is strongly recommended. Most all of the details you will require are explained in this manual. Where external devices are involved, various other manuals will benefit for the specific device. For the earlier systems, there is a large amount of documents at www.bitsavers.org. These are largely S/360 and S/370 era documents, although some later documents are available. Much more recent documents can be located at the mainframe's manufacturer's websites.

Assigned storage areas used by the hardware and sometimes SATK must be honored.

Recommended Documents

The following documents are recommended. You already have those that are supplied by SATK. This document is one of them. Bitsavers, mentioned in the previous paragraph, has many of the documents you may find helpful. To locate other documents, and sometimes the ones on Bitsavers, your Internet search engine is your friend. The easiest way to potentially find a document is to search for its document number, particularly those from the mainframe's manufacturer. Your largest selection will likely be found by excluding the version number.

SATK Related Documents

SATK User Guide – doc/guide.pdf (this document)

SATK User Guide

ASMA – A Small Mainframe Assembler – doc/asma/ASMA.pdf

Stand Alone Tool Kit Common Macros – doc/macros/SATK.pdf

Initial Program Load with ASMA – doc/asma/IPLASMA.pdf

Undocumented Hercules – doc/herc/Undocumented.pdf

Hercules Fixed Block Architecture Emulation Reference Manual –
doc/herc/FBA Manual.pdf

Bitsavers Recommended Documents

Bitsavers occasionally reorganizes their site and links change. So the interested user will need to find the document.

IBM® System/370 Principles of Operation, GA22-7000-10, September 1987.

IBM System/370 Reference Summary, GX20-1850-2, November 1976.

SATK User Guide

Program 1 – Hello World

Program 1 is the legendary “Hello World” program. It does not require use of any I/O because it uses a DIAGNOSE instruction to deliver the message, part of Hercules VM emulation.

Program: samples/guide/pgm1/pgm1.asm

Hardware Architecture: S/370

Bare-Metal Program Type: IPL Program

IPL: list-directed

Macros: None

Program Description:

```
TITLE 'PGM1 - HELLO WORLD'
```

Gives the assembly listing a title.

```
PRINT DATA
```

See all object data in listing

See all of the data generated by the assembler.

```
PSWSECT START 0,IPLPSW
```

Start the first region for the IPL PSW

* This results in IPLPSW.bin being created in the list directed IPL directory

```
PSWEC 0,0,0,0,PGMSECT
```

The IPL PSW for this program

As with all bare-metal programs, it requires an IPL PSW to be placed in absolute addresses 0-7. To accomplish this a memory region is started, named **IPLPSW**, at address 0. The PSW reflects privilege mode and the address of the program itself, **PGMSECT**. In addition to starting a new memory region, it also starts a named control section. In the list-directed IPL control file, this area of memory becomes a binary image file named **IPLPSW.bin**. Both the second operand of the **START** operation (starting the memory region) and the **PSWEC** operation (generating a PSW in the same manner as a **CCW**) are ASMA specific.

```
PGMSECT START X'300',IPLPGM1
```

Start a second region for the program itself
The IPL program is placed in a separate memory region, **IPLPGM1**. A second **START** operation is used for this. Unlike legacy assemblers, ASMA, accepts multiple **START** statements. The new control section, **PGMSECT**, has its own location counter set to absolute address **X'300'**. **X'300'** is selected to avoid hardware assigned storage areas and the SATK reserved areas in **X'200'-X'2FF'**. As with the initial **START** operation, the second one will create another file in the list-directed IPL directory, **IPLPGM1.bin**.

Both files created within the list-directed IPL directory are created from the control sections which make up each memory region. In addition, each file will be referenced, with its respective starting

SATK User Guide

location, in the list-directed IPL control file within the same directory. The ASMA -g command line option specifies the path and file name of the **control file**. The directory for each of the binary files and control file is inferred from the -g argument. The three files are placed in the same directory.

```
BALR 12,0          Establish my base register
USING *,12         Tell the assembler
```

Usual program set up for its intended base register.

```
MVC X'68'(8,0),PGMEX Set a NEW PSW to trap program interruptions
PGMEX PSWEC 0,0,2,0,X'28' Points to the program OLD PSW
```

This is the earliest point in the program that a protection against an uninterrupted program interruption loop can be established. Why? Because the MVC instruction requires a base register. Once the base register is established, the program can place into the assigned storage area used by the hardware the program interruption New PSW. This PSW gets loaded in case a program interruption occurs. It stops the CPU automatically by having the wait bit set, the 2 in the third operand. All of the maskable interrupts are disabled. The combination of the two conditions effectively stops the CPU. Hercules will display such a PSW on the console, providing an opportunity to communicate why the program came to a halt. By using X'28' as the address, the PSW indicates that it was a program interruption and the Old PSW (where the condition occurred) is at real address X'28'.

```
LM 8,10,DIAGPRMS Load the DIAGNOSE instruction's registers
DIAG 8,10,X'008' Issue the Hello World message to the console
DIAGPRMS DC A(COMMAND),A(0),X'00',AL3(CMDLEN),A(0)
COMMAND DC C'MSG * ' VM emulated command
MESSAGE DC C'Hello Bare-Metal World!' Message text sent to self
CMDLEN EQU *-COMMAND
```

This sequence initializes the registers used by the DIAG instruction, from DIAGPRMS. In turn, the DC's point to the message and contain its length, along with the flag byte.

```
BNZ BADCMD If a non-zero condition code set, the DIAGNOSE failed
LTR 10,10 Was a non-zero return code set by the DIAGNOSE?
BNZ BADCMD ..Yes, abnormal program termination
```

The results of the DIAG are checked here, first, the condition code and then the return code. If either fails the program branches to the label BADCMD. Instruction execution simply falls through if the DIAG was successful.

```
LPSW GOOD ..No, normal program termination
SPACE 1
BADCMD LPSW BAD End with a bad address in PSW field
GOOD PSWEC 0,0,2,0,0 Successful execution of the program
BAD PSWEC 0,0,2,0,X'DEAD' Unsuccessful execution of the program
```

Success or failure is indicated by which PSW is explicitly loaded by the program. In both cases a disabled wait state is established as was the case with the established Program New PSW. The difference is in the address field. For a successful program execution, an address of 0 is used. For an

SATK User Guide

abnormal termination, an address of DEAD, all hexadecimal digits, is used. The LPSW instruction loads the PSW from memory into the active PSW.

So, the instruction address of the final program PSW indicates how the program was executed as in this table.

PSW Address	Meaning
X'0000'	Successful execution of the program
X'0028'	An unexpected program interruption occurred
X'DEAD'	An abnormal program termination occurred

Step 1 – Assemble the Program

Change the working directory to the directory into which you have moved the files from SATK.

Adjust the copied asm script with local environment changes (and adjust to your scripting language if needed).

The assembly is invoked with this command in the asm script:

```
{ASMA} -t s370 -d --stats -g ldip1/pgm1.txt -l asma-$sfx.txt pgm1.asm
```

The script variable {ASMA} points to the local location of the asma.py script within SATK's tools directory.

asma.py uses a number of command-line arguments. These affect the assembly in the following ways.

- t s370 targets the S/370 architecture by the assembler. It restricts recognized instruction mnemonics to those recognized by systems implementing S/370.
- d influences the content of the listing. It causes ASMA to display in hexadecimal and characters the binary content it has created.
- stats causes asma.py to display on the console run-time time statistics and where time is spent. A statement rate is also provided.
- g ldip1/pgm1.txt identifies the type of output to be generated by the assembler, a list-directed IPL directory, and where this content is to be placed. The -d argument prints the content of the binary files.

SATK User Guide

`-l asma-${sfx}.txt` directs the assembler to generate a listing and into which file it is to be written. The use of the `$sfx` script variable allows the listing file name to include the date and time when it is created.

`pgm1.asm` identifies for the source file assembled by ASMA. It is a positional parameter and must always be the last in the command line.

Step 2 – Run the Bare-Metal Program

Change the working directory to the directory into which you have moved the files from SATK.

Execution of the program involves performing a list-directed IPL upon the `-g` file created by ASMA.

The IPL is initiated with these two commands in the `ipl` script:

```
export HERCULES_RC=pgm1.rc
${HERCULES} -v -f pgm1.conf >> log-${sfx}.txt
```

The first statement identifies to Hercules where to find the run control file, in this case `pgm1.rc`. This statement creates both a script variable and then provides it as a platform environment variable for access by Hercules. Script variables only exist within the script. Environment variables are made available to other programs.

The second statement actually launches the Hercules emulator. The `${HERCULES}` script variable identifies where on the local system the Hercules executable file is located.

The following command-line arguments influence Hercules execution.

`-v` causes Hyperion to provide verbose messages in the log.

`-f pgm1.conf` identifies to Hercules its configuration file.

The log file created by Hercules is sent to the file `log-${sfx}.txt`. As with the assembler listing, a date and time is added to the file name.

Configuration File

The Hercules configuration file has the following contents.

```
# Hercules sample configuration file for pgm1.asm
ARCHMODE S/370          # S/370 targeted by pgm1.asm
MAINSIZE 64K            # 64K minimum for S/370. Sample starts at X'300'
NUMCPU 1
DIAG8CMD enable         # Allow use of DIAG X'008'

# Console Device
000F 3215-C /           # Required by mainframe systems
```

SATK User Guide

ARCHMODE `S/370` is specified because that is the architecture targeted by the program with the assembler's `-t` command-line argument.

MAINSIZE `64K` is the minimum value supported by SDL Hercules for an S/370 system. Additionally, `64K` will probably suffice for most bare-metal programs.

NUMCPU `1` specifies one CPU which is the norm for S/370 systems. And `pgm1.asm` does not require more than one CPU.

DIAG8CMD `enable` allows the `DIAGNOSE` function `X'008'` to be used by the program. In as much as this is how `pgm1.asm` conveys the Hello World message, it is required.

The single integrated console device is specified because mainframe systems require at least one console. It is available at address `000F`.

Run Control File

The run control file is run after Hercules has configured the system with the configuration file. It contains Hercules console commands. All could be entered by means of the Hercules console itself. This simply automates the process.

The run control file has these contents:

```
# This file is intended for use with Hercules Hyperion
# Comment this statement if you do not want to trace
t+
# Perform the list-directed IPL
ipl ldip1/pgm1.txt
```

The `t+` Hercules console command causes instruction tracing to be performed. It is extremely useful during debugging of bare-metal programs, so is enabled here.

The `ipl` command performs the actual IPL, launching the bare-metal program from the list-directed IPL directory created by the `ASMA -g` argument.

List-Directed IPL Directory

The list-directed IPL directory's contents are created by the `ASMA -g` command-line argument. This directory contains three files related to the IPL function, all constructed by `ASMA`:

- `IPLPGM1.bin` – the bare-metal program's binary content as assembled (this is why the `ASMA -d` option is used).
- `IPLPSW.bin` – the IPL PSW used to start the execution of the program by Hercules following IPL.

SATK User Guide

- `pgm1.txt` – the IPL control file.

The IPL control file dictates where in memory the binary files are loaded, their starting absolute memory addresses. The address was specified in the ASMA `START` operation's first parameter. The `START` operation's second parameter specified the name of the region file.

```
IPLPSW.bin 0x0  
IPLPGM1.bin 0x300
```

When the `ipl` Hercules console command is executed by the run control file, this control file in the list-directed IPL directory controls where the file content is placed in memory.

The `IPLPSW.bin` file is placed at address `X'0'`. From the first 0 byte of storage, Hercules loads the active PSW. After which, control of the CPU is passed to the loaded bare-metal program.

The Hercules console command `quit` may be entered to terminate Hercules execution.

SATK User Guide

Program 2 – Hello World From Disk

Program: samples/guide/pgm1/pgm1.asm

Hardware Architecture: S/370

Bare-Metal Program Type: IPL Program

IPL: FBA DASD volume

Macros: None

Program Description: same as Program 1

Program 2 extends the use of Program 1. It will use the output created by Program 1 to create a FBA volume that contains Program 1, the Hello World program. Copy from the `samples/guide/pgm2` directory into **your** Program 2 directory the following files:

- `med` – which builds the IPL FBA disk
- `ipl` – which IPL's the FBA disk.

Step 1 – Assemble the Program

This was already completed in Program 1. Program 2 uses the output from Step 1 in the previous program. Simply copy the contents of the `ldipl` directory in your Program 1 directory into the directory you are using for Program 2. Copy these files into the Program 2 `ldipl` directory. Do not change the file names!!

- `IPLPGM1.bin`
- `IPLPSW.bin`
- `pgm1.txt`

Step 2 – Create IPL Medium

Creation of the FBA IPL medium is accomplished with this command:

```
${IPLASMA} -v -f ld -m pgm2.3310 --records ldipl/pgm1.txt >> iplasma-${sfx}.txt
```

This command invokes `tools/iplasma.py` to build the FBA disk. As input to the tool, it uses the copied output from ASMA in Program 1, the list-directed IPL directory. This input file is the last argument of the command: `ldipl/pgm1.txt`.

SATK User Guide

Additionally the `-f ld` argument tells `iplasma.py` the format of the input file. `ld` indicates a list-directed IPL directory control file.

The output from the tool, the emulated IPL capable FBA disk file, is specified by the `-m` argument: `-m pgm2.3310`. The `-d` or `--dtype` argument defaults to FBA which implies a 3310 device type. Other FBA device types are supported.

The output from the tool is sent to a file in your Program 2 directory:

```
iplasma-mmddyy-hhmmss.txt.
```

Two arguments influence the content of the output: `-v` and `--records`. `-v` requests verbose output. All lines lacking “`iplasma.py`” at the beginning of the output result because of the `-v` argument. When writing your own bare-metal programs, you will want this output.

The following explains the `iplasma.py` tool’s output.

```
iplasma.py Copyright (C) 2015-2020 Harold Grovesteen
```

The usual copyright notice from the author.

```
iplasma.py - IPL HWM: 00036C
iplasma.py - Used HWM: 000370
iplasma.py - IPL PSW: 0008000000000300
```

The memory high water mark of the IPL and the high water mark actually used and the IPL PSW used by the IPL are displayed.

```
Volume:  TYPE=3310 SECTORS=125664 LFS=False
Host:     FILE=64339968 (61.3MB)
Block:    LENGTH=0 SECTORS=0 BLOCKS=None
```

Description of the selected device type and its present characteristics. These characteristics relate to the general device not necessarily the device constructed by `iplasma.py`. The general 3310 volume has 125,664 sectors. `LFS` indicates whether the host requires large file system support when accessing the FBA emulated file created with these characteristics. The `Host` line describes the full file size of the emulating file. The `Block` line indicates the file length, number of sectors and blocks currently created. These values should always be 0 without any blocks because these values are provided prior to the volumes actual creation.

SATK User Guide

FBA DASD Map:
IPL0: 0-0
VOLLBL: 1-1
IPLPGM1.bin: 2-2

How the FBA volume has been allocated for the IPL usage is described by these output statements. The range indicates the beginning and ending sectors for each item. In this case, only one sector has been reserved for each major FBA disk items.

3310 IPL: LDIPL

This identifies the type of source installed on the FBA volume.

```
iplasma.py - IPL Record 1:
000588 43000598 40000008 42000300 00000200
000598 06000001 00000002
iplasma.py - IPL Record 0 - CCWs:
000008 02000570 40000200 08000588 00000001
iplasma.py - IPL Record 0:
000000 00080000 00000300 02000570 40000200
000010 08000588 00000001
```

These statements identify where in memory (and the binary content for logical IPL records 0 and 1) and the order created by the tool. For IPL Record 0's creation, you must know where IPL Record 1 is placed in memory. Hence, IPL Record 1 is built before IPL Record 0. This is illustrated by the addresses in bold, **X'000588'**. The first address reflects where IPL Record 1 will be placed. This can then be used in IPL Record 0's creation of the Transfer-In-Channel CCW's address that causes the CCW program chain to continue with the CCW's in IPL Record 1. The following Memory Map is also required to know where these components reside.

Memory Map:
IPL0: 000000-000017
PSW: 000000-000007
CCW0: 000008-000017
IPLPGM1.bin: 000300-00036C
IPL1: 000588-00059F

Memory is allocated for the IPL process as described by these output lines. This is where you should find the specified content in memory. Addresses are in hexadecimal. Useful for debugging.

```
iplasma.py - emulated medium pgm2.3310 created with file size: 1536
```

The file created and its physical size in the host's file system is identified here. Note that the physical size is significantly smaller than that identified in the above Host output line for a generic FBA 3310 device.

This results from the `--size` argument defaulting to `mini`. Hercules supports non-standard device sizes. For such devices the device type of the configuration file dictates the characteristics of the device, not the file size. The underlying tools used by SATK for DASD creation allows for three size

SATK User Guide

options: `mini` (the default used here), the absolute minimum file size required by the bare-metal program and supported by Hercules; `comp`, compression **eligible** file; and `std`, the normal full size volume. SATK does not create compressed files. Either `comp` or `std` create uncompressed files that may be compressed by use of a Hercules utility. Emulating files created using `mini`, as in this case, can not be compressed for Hercules use.

```
FBA sector 0
 000000 00080000 00000300 02000570 40000200
 000010 08000588 00000001 43000598 40000008
 000020 42000300 00000200 06000001 00000002
 000030 00000000 00000000 00000000 00000000
...
 0001F0 00000000 00000000 00000000 00000000
FBA sector 2
 000000 05C0D207 0068C026 988AC03E 838A0008
 000010 4770C01C 12AA4770 C01C8200 C02E8200
 000020 C0360000 00000000 000A0000 00000028
 000030 000A0000 00000000 000A0000 0000DEAD
 000040 00000350 00000000 0000001D 00000000
 000050 D4E2C740 5C40C885 93939640 C2819985
 000060 60D485A3 819340E6 96999384 5A000000
 000070 00000000 00000000 00000000 00000000
...
 0001F0 00000000 00000000 00000000 00000000
```

These lines result from the `--records` command-line argument. The final output is that of each written sector on the FBA volume by `iplasma.py`. The redundant lines of binary zeros are indicated by the `...` in the above lines. The “address” on the left indicates the displacement into the sector where the data was placed. Sector 1 is not shown because the volume label has not been written to the emulated disk and hence is excluded from this text output. However, sector 1 is all zeros in the emulated disk.

For this disk, three sectors exist. 512 bytes per sector results in a file of 1536 bytes.

SATK User Guide

Step 3 – Run the Bare-Metal Program

The Hercules `pgm2.conf` file has added one line.

```
# Hercules sample configuration file for pgm2.asm
ARCHMODE S/370          # S/370 targeted by pgm2.asm
MAINSIZE 64K            # 64K minimum for S/370. Sample starts at X'300'
NUMCPU 1
DIAG8CMD enable         # Allow use of DIAG X'008'

# Devices
000F 3215-C /           # Console required by mainframe systems
0110 3310 pgm2.3310     # FBA IPL disk
```

As can be seen, the FBA disk created by Step 2 has been added to the configuration file as device `110` (in hexadecimal). Otherwise it is identical to the configuration file used by Program 1.

Likewise the run control file has been changed to cause an IPL from the FBA disk created in Step 2.

```
# This file is intended for use with Hercules Hyperion
# Comment this statement if you do not want to trace
t+
t+110
# Perform the FBA based IPL
ipl 110
```

The additional `t+` line with the FBA device address, `t+110`, causes more tracing. This adds tracing of the operation of the CCW's during the IPL process. Again, for debugging you may want this.

SATK User Guide

Program 3 – Hello World Using Console I/O

Program: samples/guide/pgm1/pgm3.asm

Hardware Architecture: S/370

Bare-Metal Program Type: IPL Program

IPL: FBA DASD volume

Macros: ARCHLVL, ASAIPL, ASALOAD, ASAREA, DWAIT, DWAITEND

References:

IBM System/370 Reference Summary

doc/asma/IPLASMA.pdf

doc/macros/SATK.pdf

Program Description:

Program 3 is another Hello World program, but what changes is how the output is delivered to the Hercules console. This program uses the console device to display the message using hardware input/output commands (as opposed to a **DIAG** instruction).

S/370 PoO, Chapter 13, Input/Output Operations, is recommended for details. “Control of Input/Output Devices” describes the various privileged instructions used by a bare-metal program. “Execution of Input/Output Operations” describes how the channel-command words are used to actually drive operations with a device. “Conclusion of Input/Output Operations” discusses how such operations end, the Input/Output Interruption, and the Channel-Status Word (CSW) which describes the conditions related to the I/O operation.

Assigned Storage Areas

By intent, this program uses few pre-defined storage structures. They tend to hide the basic actions of the code which is the objective of this program. However, DSECT's are valuable. The DSECT macro, a convenience for defining various DSECT definitions within a program, can generate a number of used DSECT's. The format is either specified by the hardware or how SATK, software, uses the structure.

The hardware defines the usage of various areas of storage, generally referred to as Assigned Storage Areas. Most of these areas are within the first 512 bytes of storage, although additional areas exist in

SATK User Guide

the second 4K-byte page. Two macros are provided by SATK for definition of these areas: ASAREA and ASAZAREA, respectively. The ASAZAREA macro applies only to z/Architecture® assigned areas and is not discussed here.

ASA ASAREA DSECT=YES

The ASAREA macro defines the areas assigned by hardware between X'0' and X'1FF' regardless of architecture. In general, the area between X'200'-X'11FF' is available for program usage. In the context of hardware specified memory usage, SATK is the “program.” SATK defines how the storage from X'200'-X'11FF' is used. The following table identifies each real memory area used by SATK. EOM stands for “end of memory”. The **red** areas are assigned by hardware. The **blue** area is required when a Boot Loader is in use. The **gray** area is reserved by SATK for future usage. All of the white areas are optional as needed by the program.

Area	Start	End	Length	Description
A	000	1FF	512	Assigned by hardware
B	200	23F	64	Hercules IPL command parameters
C	240	28F	80	LOD1 IPL Record 4
D	290	2FF	112	Reserved by SATK
E	300	11FF	3840	Small IPL program or Boot Loader
F	1200	1FFF	3584	Assigned by hardware
G	2000	EOM	varies	Any IPL program or Booted program

The bare-metal program will interact with any of these areas. How each area or field is used by the assigner should be respected by the program.

ARCHLVL

Architecture level is an SATK concept. Do you care about this? Not in this program. It is specific to S/370 using Extended Control mode. So why is it there? Many SATK macros **are** architecture sensitive. They generate different things based upon the architecture level of the program.

An MNOTE message is always produced by ARCHLVL informing the user which architecture has been identified. In this case the architecture level is 3. 3 stands for S/370. So, the correct target (ASMA command-line argument - t) has been specified when the assembler was initiated by the program.

The following two macros are sensitive to the architecture level of the program.

ASASECT ASALOAD
ASAIPL IA=PGMSTART Define the bare-metal program's IPL PSW

SATK User Guide

These two macros fill in the ASA with interrupt trap PSW's (ASALOAD) and an IPL PSW (ASAIPL). ASALOAD creates the ASA region image placed within the list-directed IPL directory, naming it the default ASAREGN. Note that both the region and control sections created by ASALOAD are symbols within the assembler's symbol dictionary. They must be unique.

CPU Initialization

PGMSECT START X'300',IPLPGM3 Start a second region for the program itself
The IPL Program itself has its own memory region and control section, IPLPGM3, and PGMSECT, respectively. Its load address is X'300'

USING ASA,0 Give me instruction access to the ASA CSECT
Because the program uses the ASA fields as defined by the DSECT created with the ASAREA macro, a USING operation is required to establish its addressability. Because it is within the first 4K bytes of memory, base displacement addressing does not require an active base register. Register 0 (ignored by such instructions) acts as a "filler". The assembler does not know the difference.

PGMSTART BALR 12,0 Establish my base register
USING *,12 Tell the assembler

The usual program initialization of its (real) base register.

MVC RSTNPSW,PGMRS

The ASA area for a console operator restart of the program is at real address X'0'-X'7'. Prior to the IPL these same addresses are assigned to the IPL PSW. I/O commands can only access absolute storage. However, to support multiprocessing, the CPU can map the first 4K, or 8K when using z/Architecture, bytes of storage to different absolute locations by means of the prefix register. Real storage, as seen by the CPU can be different than the same absolute addresses. Real storage then is a mechanism exclusive to the CPU. I/O operations know nothing about it.

The system reset that occurs at the start of the IPL function, will set the prefix register to 0. So, for a bare-metal program immediate following IPL, the locations X'0'-X'7' refer to the identical physical storage addresses for both modes of memory access: absolute (the I/O of the IPL) and real (instruction execution) storage.

So why does the bare-metal program care? Maybe it doesn't. But Program 3 is not designed to be reentered by an operator initiated restart. It is not serially reusable. And these 8 bytes, as real addresses, still contain the IPL PSW (as I/O data from the IPL). Left alone, the operator could initiate a restart using (now real addresses) the IPL PSW, entering the program as if it had just been IPL'd.

To ensure this does not happen, the Restart New PSW, at real addresses X'0'-X'7' are initialized after IPL with a trap PSW. Because these same addresses are used for the IPL PSW, this change can not occur until **after** the IPL.

SATK User Guide

Input/Output Initialization

	LCTL	2,2,CR2	Enable only channel 0 for I/O interruptions
CR2	DC	XL4'80000000'	Enable only channel 0 where console is connected

When the system reset occurs at the start of the IPL function, control register 2 is initialized to allow interruptions to be presented to the CPU, provided the PSW allows them, for each channel attached to the CPU. This would be channel 0, for the console device (00F), and channel 1, for the FBA IPL disk device (110). In other words, CR2 would contain X'C0000000'. By loading control register 2, used for controlling interruptions to the CPU from various channels, with X'80000000', the FBA disk can not interrupt the CPU even if an interruption were to occur. If it were to occur, the program would fail with an error condition. By loading control register 2 with this value, makes the program a bit more fail proof. This takes care of the channels.

	LH	1,CONDEV	Set up I/O device address in I/O register
	TIO	0(1)	Determine if the console is there
CONDEV	DC	XL2'000F'	Console device address

Now, the subchannel, well, really, the channel, subchannel (control unit) and device are inspected to determine if they are operational, the complete path is functional or not.

The TIO (TEST IO) instruction sends a special command to the device, one not possible via a CCW. This command is X'00', a value invalid for the command field of a CCW. The command is treated as an immediate command (no data transferred) by the control unit. How the command succeeds or fails will indicate the state of the console device's path and will, accordingly, influence the return code set in the PSW by the TIO instruction.

The TIO instruction uses the storage format. The first (and only) operand identifies the channel/unit address of the targeted device. The storage format of the instruction, typically allows two ways that this address is supplied:

- as a hard coded address, for example, TIO X'00F' (using just the displacement field of the instruction), or
- as a generic value placed in a register, for example TIO 0(1).

Program 3 uses the register approach. Before executing the TIO instruction, general register 1 is loaded with the channel and unit address from a field initialized at assembly time, CONDEV.

The TIO instruction attempts to send the command, restricted for its use, to the device identified in the instruction's operand. X'00F' in this case.

BC	B'0100',DEVNOAVL	..No, CC=1 might have a different address
BC	B'0010',DEVBUSY	..No, CC=2 console device or channel is busy
BC	B'0001',DEVCSW	..No, CC=3 CSW stored in ASA at X'40'

SATK User Guide

Depending upon the condition code setting, different error cases are identified. If a branch is taken it passes control to a LPSW that establishes a terminating error condition.

If none of the three branches is taken, then the console device is available on the path specified by the device address. The final step in initializing the input/output operation can proceed.

```

                MVC    CAW(8),CCWADDR    Identify in ASA where first CCW resides
CCWADDR        DC     A(CONCCW) Address of first CCW executed by console device.
CONCCW         CCW0   X'09',MESSAGE,0,MSSGLEN    Write Hello World message with CR
*              CCW0   X'03',0,0,1                ..then a NOP.
* If the preceding NOP CCW command is enabled, then the CONCCW must set
* command chaining in the flag byte, setting the third operand to X'40'
MESSAGE        DC     C'Hello Bare-Metal World!' Data sent to console device
MSSGLEN        EQU    *-MESSAGE                Length of Hello World text data
```

A field within the ASA is reserved to identify where the CCW chain resides for execution by the channel, the channel address word (CAW). In this case, an address field is moved to the ASA, CCWADDR. In turn, this field points to the location of the CCW chain at CONCCW.

The CCW chain consists of a single CCW, a Write with CR (carriage return), command code X'09'. The CCW points to the data being sent to the console, namely, the Hello World message at MESSAGE, and its length, MSSGLEN. The only instruction here moves an address to the CAW. All the other assembler statements provide the initialization for the input/output operation itself.

Why is a NOP CCW, command X'03', commented out? On good authority, while unlikely when working with Hercules, in the real world it is possible for the ending status to be delayed or missed. By chaining the NOP CCW to the Write with CR command, ending status is sure to be presented to the application.

Everything is now ready to output the Hello World message.

Input/Output Operation

Everything is now ready to actually send the Hello World message to the console. This is accomplished by the SIO (Start I/O) instruction.

```

SIO    0(1)          Request console channel program to start, did it?
BC     B'0100',DEVNOAVL ..No, CC=1 don't know why, but tell someone.
BC     B'0010',DEVBUSY  ..No, CC=2 console device or channel is busy
BC     B'0001',DEVCSW   ..No, CC=3 CSW stored in ASA at X'40'
```

The SIO is similar in design to the TIO previously used. It requires the device address, still in general register 1 from the TIO. It too sets the condition code based upon the results of the operation at the device when the CCW command is sent to the device. Unlike the TIO that sends a specific command to the device, SIO sends the command as specified in the first CCW. Following the reception of the command by the device, an initial status is sent to the channel attached to the CPU. As with the

SATK User Guide

immediate command sent by the TIO instruction, the SIO will set the condition code based upon the device's response to the command.

The branch instructions detect anything that represents a failure, and, as with the TIO, transfers control to one of these locations in the program. They each load their respective PSW's and abnormally terminate the program with the designated instruction address.

```
DEVNOAVL LPSW  NODEV      Code 004 End console device is not available
DEVBUSY  LPSW  BUSYDEV    Code 008 End because device is busy (no wait)
DEVCSW   LPSW  CSWSTR     Code 00C End because CSW stored in ASA
NODEV    DWAIT PGM=03,CMP=0,CODE=004 Console device not available
BUSYDEV  DWAIT PGM=03,CMP=0,CODE=008 Console device busy
CSWSTR   DWAIT PGM=03,CMP=0,CODE=00C CSW stored in ASA
```

By the process of elimination, success is recognized and the CPU continues with the next sequential instruction. At the device, the command has been accepted and the transfer of the Hello World message to the device is occurring. Hence this comment in the program.

* Console device is now receiving the message (CC=0)

Input/Output Interruption

The transfer to the device of the message continues until it is done some time in the future. What is the program to do during this time? Wait for the transfer to complete. Some discussion of program interruptions in general is useful at this point.

The mainframe has six interrupt classes. Most provide some information about the interruption's cause. Some can be masked off or on. If an interruption class is masked off and an interruption of that class occurs, it will remain pending until such time as the class is masked on. Other interruptions are unmasked and are always enabled. The following table summarizes the general aspects of each class. See Chapter 6 in the S/370 PoO for details. Addresses are real within the ASA.

Priority	Class	Class Mask	Interrupt Mask	Old PSW	New PSW	Interruption Information
0	Machine Check - Exigent	None	None	X'30'	X'70'	None
1	Supervisor Call	None	None	X'20'	X'60'	X'88'-X'8B'
2	Program	None	None	X'28'	X'68'	X'8C'-X'8F'
3	Machine Check - Repressible	PSW	CR 14	X'30'	X'70'	X'E8'-X'EF'
4	External	PSW	CR 0	X'18'	X'58'	X'84'-X'87'
5	Input/Output	PSW	CR 2	X'38'	X'78'	X'BA'-X'BB', CSW X'40'-X'47'
6	Restart	None	None	X'08'	X'00'	None

SATK User Guide

Regardless of the class, an interruption stores the active PSW and any related interruption information into the class' ASA old PSW area and its interruption information area. This is followed by loading the active PSW from the class' ASA new PSW area. Before an interruption occurs or can occur due to masking, the class' New PSW must be set with meaningful PSW content. If this has not occurred, results are technically unpredictable, but usually result in an interrupt loop. Very unpleasant! The trap PSW's provide a valid PSW and will disable all maskable interruption classes if loaded by an unexpected interruption.

As seen previously, CR 2 was initialized to allow only input/output interruptions from channel 0. However, the I/O new PSW currently has only the trap PSW. The active PSW was loaded by the IPL function and it has input/output interruptions disabled.

The next instruction addresses the issue concerning the I/O new PSW.

```
DOWAIT   MVC   IONPSW(8),CONT   Set PSW for after I/O interrupt
          LPSW  WAIT           Wait for I/O interruption and CSW from channel
IODONE   EQU   *               The bare-metal program continues here after I/O
          MVC   IONPSW(8),IOTRAP   Restore I/O trap PSW
CONT     PSWEC 0,0,0,0,IODONE   Causes the CPU to continue after waiting
WAIT     PSWEC 2,0,2,0,0       Causes CPU to wait for I/O interruption
```

It copies a different PSW, at label CONT, into the I/O New PSW ASA. This PSW, like the IPL PSW, disables maskable interruptions, including any input/output interruptions, and does not cause the CPU to wait. In other words, when loaded by an interruption, the CPU begins executing instructions at the location IODONE.

The instruction following the MVC, the LPSW actually changes the active PSW to become a waiting PSW. The CPU ceases executing instructions. However, in addition to waiting, this active PSW has enabled I/O interruptions. These two instructions and two PSW's work in conjunction to cause the wait for the I/O interrupt and continue CPU instructions after the interruption has occurred. The very next thing that occurs is restoring the trap PSW for I/O interruptions in case another one occurs.

Input/Output Completion Analysis

Now that the program knows the I/O is completed by the device, how well did it go? First, because there is no control over which device might present an interruption when enabled, was it the console that provided the interruption?

```
CH       1,I0ICODE           Is the interrupt from the console?
BNE      DEVUNKN             ..No, end program with an error
DEVUNKN  LPSW  NOTCON   Code 010 End unexpected device caused I/O interruption
NOTCON   DWAIT  PGM=03,CMP=0,CODE=010   interruption from some other device
```


SATK User Guide

General register 1 still has the console device address, so it is used to examine the I/O interruption code, the device address that triggered the interruption. If it is not the console device the program loads a disabled wait PSW to indicate a device unrecognized by the program caused an interruption.

Now that it is known that the interruption came from the console, the I/O operation's status can be checked in the Channel Status Word (CSW).

OC	STATUS,CSW+4	Accumulate Device and Channel status
STATUS DC	XL2'0000'	Used to accumulate unit and channel status

One of the realities that the program must address is that the ending status of the channel and the device are usually presented to the program at the same time, this does not have to be the case and in fact some device operations are designed to present separate status interruptions. To handle this situation the program combines the status of multiple interruptions to determine the overall status of the channel and device. The OC does that by combining the current known status, in this case initially all zeros, with the status presented by the interruption. Chapter 16, pages 13-62 – 13-78, of the S/370 PoO details the content of the CSW and the various status conditions. The combined status is inspected for any error conditions, as follows.

CLI	STATUS+1,X'00'	Did the channel have, a problem?
BNE	CHNLERR	..Yes, end with a channel error

The first check is for channel level errors. Many of the detected conditions are programming created error conditions. These are useful in debugging. Others relate to physical channel issues, none of which can happen with use of an emulator, such as Hercules. Nevertheless, because SATK strives to function in all mainframe contexts, the checks are still performed.

The next set of status conditions are reported as unit status. Some may be errors, some not depending on the device. For the console, all of these are considered as errors.

TM	STATUS,X'F3'	Did the unit encounter a problem?
BNZ	UNITERR	..No, end with a unit error

If none of these conditions happened, then all that is left is to determine if the I/O has completed. This requires both the Channel End and Unit End conditions to be set.

TM	STATUS,X'0C'	Did both channel and unit end?
BNO	DOWAIT	Wait again for both to be done

The test assumes that if either status (or both) have not been presented, they eventually will. The CPU is then caused to wait again by branching back to the DOWAIT label, and wait for an interruption.

When both the Channel End and Unit End conditions have been signaled, the branch is ignored and the program falls through to the next instruction. The message has been successfully sent to the console.

LPSW	DONE	Normal program termination
DONE	DWAITEND	Successful execution of the program

And that is it! The program is done. The PSW loaded has all PSW interrupt masks off and waits, a condition from which the CPU will not exit (except by a manually triggered restart, which of course

SATK User Guide

was earlier trapped). The DWAITEND macro creates the appropriate architecture level PSW for the program to enter this termination state.

Step 1 – Assemble the Program

This program utilizes some macros supplied with the SATK in a macro library. Macro libraries are accessed via the environment variable MACLIB. This environment variable requires the SATK delivered macro library: `${SATK}/mac lib`.

```
export MACLIB=${REPO}/mac lib      # Access SATK supplied macros
${ASMA} -t s370 -d --stats -g ldip l/pgm3.txt -l asma-$sfx.txt pgm3.asm
```

The above statements are from the `asm` script in the `pgm3` directory. All that has really changed from the previous programs is the addition of the `export` statement that creates the `MACLIB` environment variable.

Step 2 – Create the IPL Medium

Again the `med` script is used to create the IPL medium for Program 3.

```
${IPLASMA} -v -f ld -m pgm3.3310 --records --asa=ASAREGN.bin \
ldip l/pgm3.txt >> ip lasma-${sfx}.txt
```

The `ip lasma.py` command line is very similar to the one provided for Program 2. However, a new argument is provided: `--asa=ASAREGN.bin`. This argument informs `ip lasma.py` that a PSW region has been replaced by an assigned storage area region. The ASA region's file name, including its extension, is provided. This argument implies that the IPL PSW is at the beginning of the region. The other arguments are the same (other than for the adjustment to the "pgm3" filenames).

The FBA sector allocation map has changed as a result of including the `--asa` argument.

```
FBA DASD Map:
IPL0: 0-0
VOLLBL: 1-1
ASA: 2-2
IPLPGM3.bin: 3-3
```

A new sector has been allocated for the ASA area. Why? Does it not contain the IPL PSW? Yes. But, the FBA sector zero must contain both IPL records 0 **and** 1. Record 1 contains the rest of the channel program that performs the IPL. So, the ASA content is moved to the third sector and the channel program is adjusted to read this sector. Where did the IPL PSW come from? It was extracted from the ASA binary image file, bytes 0-7, and placed at the start of IPL record 0 as required by the IPL process.

```
Memory Map:
IPL0: 000000-000017
ASA: 000000-0001FF
CCW0: 000008-000017
```

SATK User Guide

IPLPGM3.bin: 000300-00040F

IPL1: 000628-000657

As can be seen from the memory map, the ASA region has been allocated to the first 512 bytes of memory, as is expected from the assembly. The areas in **bold** have been created by `iplasma.py` from its knowledge of the intent of the ASALOAD and ASAIPL macros. Bytes X'0'-X'17' constitute IPL Record 0 (sector 0 content in **bold**). Bytes X'18'-X'48' contain IPL Record 1 (sector 0 content in *italics*).

```
FBA sector 0
000000  00080000 00000300 02000610 40000200
000010  08000628 00000001 43000648 40000008
000020  42000300 40000200 43000650 40000008
000030  42000000 00000200 06000001 00000003
000040  06000001 00000002 00000000 00000000
...
0001F0  00000000 00000000 00000000 00000000
FBA sector 2
000000  00080000 00000300 00000000 00000000
000010  00000000 00000000 00000000 00000000
000020  00000000 00000000 00000000 00000000
000030  00000000 00000000 00000000 00000000
000040  00000000 00000000 00000000 00000000
000050  00000000 00000000 000A0000 00000018
000060  000A0000 00000020 000A0000 00000028
000070  000A0000 00000030 000A0000 00000038
...
0001F0  00000000 00000000 00000000 00000000
FBA sector 3
000000  05C0D207 0000C0B6 B722C086 4810C08E
000010  9D001000 4740C06E 4720C072 4710C076
000020  D2070048 C08A9C00 10004740 C06E4720
000030  C0724710 C076D207 0078C0C6 8200C0BE
000040  D2070078 C0CE4910 00BA4770 C07AD601
000050  C0900044 9500C091 4770C07E 91F3C090
000060  4770C082 910CC090 47E0C034 8200C0D6
000070  8200C0DE 8200C0E6 8200C0EE 8200C0F6
000080  8200C0FE 8200C106 80000000 00000398
000090  000F0000 00000000 090003A0 00000017
0000A0  C8859393 9640C281 998560D4 85A38193
0000B0  40E69699 93845A00 000A0000 00010008
0000C0  020A0000 00000000 00080000 00000340
0000D0  000A0000 00000038 000A0000 00000000
0000E0  000A0000 00030004 000A0000 00030008
0000F0  000A0000 0003000C 000A0000 00030010
000100  000A0000 00030014 000A0000 00030018
000110  00000000 00000000 00000000 00000000
...
0001F0  00000000 00000000 00000000 00000000
```

Sector 1, as before, is omitted. A volume identification record has not been supplied. Sector 2 has the full contents of the ASA.

SATK User Guide

FBA IPL Channel Program

The IPL **channel program** in Sector 0 does the following:

1. The **first 24 bytes** of Sector 0 are read into storage at absolute address 0 by the hardware. The IPL function does this by sending a Read IPL command, X'02', to the IPL device. An **implied** CCW can be thought of as residing at absolute address 0. This implied CCW is not physically placed at address 0, but the hardware functions as if it were. The contents of this implied CCW causes the channel hardware to function as if the CCW was:

02000000 60000018, or in assembler format,
CCW X'02', 0, B'01100000', 24.

This implied CCW, a Read IPL command, reads 24 bytes starting at absolute address 0, with the command chain and suppress any incorrect length indication flags set. Command chaining is required to cause the IPL function's channel commands to continue with the CCW at address X'8'. The suppress incorrect length indication is required because Sector 0 is actually 512 bytes in length, but the implied CCW is only going to read the first 24 of those 512 bytes. As a result, an incorrect length would be indicated, but is suppressed by the implied flag.

2. The channel program continues, and with it the IPL function, with the CCW at absolute address 8. This CCW re-reads **all 512 bytes** of Sector 0. The first sector is read into memory addresses X'610'-X'80F'. Now a complete copy of IPL Records 0 and 1 reside in memory. IPL Record 0, in regular text, is repeated at X'610'-X'627'. IPL Record 1 now resides at X'628'-X'658'.

000610	00080000	00000300	02000610	40000200
000620	08000628	00000001	43000648	40000008
000630	42000300	40000200	<u>43000650</u>	<u>40000008</u>
000640	<u>42000000</u>	<u>00000200</u>	06000001	00000003
000650	<u>06000001</u>	<u>00000002</u>	00000000	00000000

3. The last thing that IPL Record 0 does is pass control of the CCW chain to address X'628', this time ignoring IPL Record 0, and continuing with IPL Record 1 in the second complete Sector 0.
4. The IPL program is then read into memory by the CCW's at X'628'-X'637' (in *italics*). These two CCW's use the information at X'648'-X'64F' to identify the starting sector, 3, and the number of sectors, 1, that constitute the content of the program.
5. The last thing done in the IPL Record 1's continuation of the channel program is the reading of the ASA starting at absolute address 0. The underlined CCW's at X'638'-X'657' use the information at addresses X'650'-X'657' to read 1 sector starting with sector 2. Sector 2 is the ASA region. This wipes out the IPL PSW and the two CCW's of IPL Record 0 read into memory at the beginning of the IPL function.

SATK User Guide

6. At the conclusion of the IPL channel program CCW's, the IPL device address will be stored in the I/O Interrupt ID field at absolute X'BA'. A CSW will also be stored at X'40'. If an I/O error occurred, the IPL function ceases at this point and will not load the IPL PSW.
7. A successful IPL completes with the loading of the IPL PSW (now from the ASA region) and turning control over to the CPU.

Step 3 – Run the Bare-Metal Program

The `ipl` script will, as before, executes the program. This script has been adjusted for `pgm3` file names.

The Hercules configuration file contains these statements:

```
# Hercules sample configuration file for pgm3.asm
ARCHMODE S/370          # S/370 targeted by pgm3.asm
MAINSIZE 64K             # 64K minimum for S/370. Sample starts at X'300'
NUMCPU 1
DIAG8CMD enable         # Allow use of DIAG X'008'

# Devices
000F 3215-C /           # Console required by mainframe systems
0110 3310 pgm3.3310     # FBA IPL disk
```

The only change has been to change the file name of the emulated FBA disk to Program 3.

The Hercules run control file, `pgm3.rc`, is as follows.

```
# This file is intended for use with Hercules Hyperion
# Comment either or both statements if you do not want to trace
t+
t+00F
# Perform the FBA based IPL
ipl 110
```

A single statement has been added since Program 2: `t+00F`. This statement causes tracing of the device, the console. In the Hercules log we can find this:

```
23:04:51 /Hello Bare-Metal World!
```

The `/` tells us that the content of the line came from the bare-metal program (and not Hercules) on the integrated console. There is our “Hello World” using mainframe channel I/O!

SATK User Guide

Program 4 – A Booted Hello World Program

Program: samples/guide/pgm4/pgm4.asm

Hardware Architecture: S/370

Bare-Metal Program Type: Booted Program

IPL: FBA DASD volume

Macros: ARCHLVL, ASAIPL, ASALOAD, ASAREA, DWAIT, DWAITEND

Program Description: same as Program 3

Again we announce to the world we are here. But this time the Hello World program uses a boot loader that loads it into memory. The boot loader becomes the program that uses the IPL function to bring it into memory and start its execution. In turn, the boot loader brings the primary bare-metal program, the Hello World program, into memory and passes control to it.

Copy the following files to your work directory:

- `pgm4.asm` – the booted program (Hello World) source
- `boot4.asm` – the FBA boot loader source
- `asm` – the script that assembles the booted program
- `basml` – the script that assembles the boot loader
- `med` – the script that creates the FBA emulated volume from the two assemblies
- `pgm4.conf` – the Hercules configuration file for `pgm4`
- `pgm4.rc` – the Hercules run control file for `pgm4`
- `ipl` – the script that runs the boot loader and booted program in Hercules

Step 1a – Assemble the Booted Program

Use the script `asm` to assemble program 4. This is the same program as used in Program 3, with two exceptions. As the booted program, its residence is specified by the `iplasma.odt` or `iplasma.pdf` manual as X'2000'.

The statement starting the region containing the booted program is changed for `pgm4.asm` to

```
PGMSECT  START X'2000',BOOTED4 Start a second region for the program itself
```

SATK User Guide

This changes the program's starting location to X'2000', and creates the program content file `BOOTED4.bin`.

The other exception relates to the disabled wait states generated by the Hello World program. These have been changed to reflect a component of 1. This lets the user know which program, the boot loader or the booted program, generated the error condition.

```
NODEV      DWAIT PGM=04,CMP=1,CODE=004 Console device not available
```

This causes a disabled wait state address to be X'041004'.

Step 1b – Assemble the Boot Loader

A new script was introduced allowing the assembly of the boot loader used by the Hello World: `basm`. The source of the boot loader is in `boot4.asm`.

In very simple terms, the boot loader after initialization, uses a loop to read the directed records into storage and move the contents of each into its respective memory location. Once all of the records have been read, the boot loader branches to the booted program, passing control of the CPU and environment to it. That's it.

The boot loader relies upon the LOD1 Record, IPL Record 4, for its successful execution. The structure of this SATK specific logical IPL record is documented in the `iplasma.pdf` or `iplasma.odt` files. This boot loader is somewhat of a minimalist approach to boot loaders. It does not have all of the bells and whistles that might exist in a boot loader. It **is** independent of the booted program itself. Meaning that this boot loader could be used with other programs. Any linkage to the booted program is provided by `iplasma.py` and resides on the IPL medium as the LOD1 record. So, if you need a boot loader, have at it.

How Fixed-Block Architecture I/O Works

To understand how the boot loader functions, it is necessary to understand some of how FBA DASD's operate. FBA DASD is an enhancement on top of CKD DASD. An FBA device is really, under the covers, just another CKD device. It has cylinders and tracks just like a CKD device. However, externally, those are hidden from the user. On each track there are a set of records just like a CKD device. However, unlike CKD, the records are all of a fixed length, 512 bytes. On FBA devices, records are called sectors. Outside of the hardware, a user "sees" only sectors. Physical sectors (physical records) are identified by a number, starting at zero. This section will be enhanced by use of the `FBA Manual.odt` or `FBA Manual.pdf` in the `doc/herc` SATK directory.

The context into which a FBA device normally resides is the same as that for a CKD device. Users run programs that access datasets on the device. Where a dataset resides on the volume is defined by a

SATK User Guide

Volume Table of Contents (VTOC). FBA devices usually have one of those, just like a CKD device. Each allocated dataset has an entry in the VTOC that identifies where on the volume the dataset resides.

Built into the device is a protection mechanism ensuring a program does not corrupt the data in other datasets. It does this by defining an “extent”. This CCW command, required for each sequence of channel commands, identifies to which physical sectors this sequence is allowed access. The device assumes an operating system will ensure this command is attached to each CCW sequence providing access to a dataset. Three key pieces of information are provided for this protection:

- the starting physical sector of the dataset extent,
- the first sector (relative to the start of the dataset) allowed access, and
- the last sector (relative to the start of the dataset) allowed access.

The DEFINE EXTENT CCW is about restricting a program’s access to just the dataset to which it has access. For a specific dataset, the information in this command will be constant.

The internal seeking to a specific cylinder and track occurs via the LOCATE command. It is tied to the DEFINE EXTENT by means of the dataset relative sectors. The first sector within the dataset being accessed is identified along with the number of sectors being read or written. Which is the case is indicated by the “operation code” of the command’s data. In most cases, a program will read or write the same number of sectors (bytes really, but this command only knows about sectors). The number of sectors tends to be constant. It is the starting sector number that will increment for sequential access.

Finally, the program will continue with a READ or WRITE CCW command that accesses the amount of data expected by the program.

The only exception to this pattern, is the IPL function. A DEFINE EXTENT command is not required for the IPL process because the IPL READ command will automatically define an extent that is the entire FBA volume. Succeeding LOCATE commands will therefore refer to physical sectors rather than sectors relative to the dataset. In this special case, a physical sector and dataset sector are the same.

Now enters bare-metal programming. There is no VTOC on the device. There could be of course, but at present it does not exist. All that exists are the directed records that the boot loader expects to read. It is possible to determine the size of the volume using the READ DEVICE CHARACTERISTICS command. But that is extra work. Each access by the boot loader of a directed record is going to require each of the three commands described above. Instead of trying to figure out an extent for the entire volume (as is done during the IPL function), the boot loader will change the extent so that just the number of sectors being read will be defined. With this strategy, rather than the LOCATE

SATK User Guide

command reflecting the movement of reads through the dataset, the extent will march through the directed records as each is brought into memory. Only one field of information for the read operation requires changing as each directed record is read: the physical sector starting the extent. All other fields will remain constant.

The LOD1 record in storage contains the critical values for reading the directed records:

- at address X'246' (LOD1MDLN), the maximum directed record size (in bytes),
- at address X'258' (LOD1FSEC), the starting physical sector of the directed records, and
- at address X'26C' (LOD1IOA), the boot loader I/O area (where directed records are read).

The following instructions from `boot4.asm` and storage locations initialize these I/O related locations:

```

        ICM    6,B'0011',LOD1MDLN    Fetch from LOD1 the maximum directed length
        STH    6,FBACW3+6    CCW - Will always read the same number of bytes
        SRL    6,9(0)        Convert bytes into sectors.
        STH    6,LOCSECS      LOC - Will always read the same number of sectors
        LR     5,6            Save the number of sectors (updates extent)
        BCTR   6,0            And the logical end of the extent is constant
        ST     6,ENDLSEC      EXT - Set it in the FBA extent data
* R6 is now available for other uses
        L      4,LOD1IOA      Remeber where to find a directed record
        STCM   4,B'0111',FBACW3+1    CCW - Where the loader read's its data
        MVC    FRSTPSEC,LOD1FSEC      EXT - Set starting sector of first record
        ...
FBACW1   CCW0   DEFN_EXT,EXTENT,CC,EXTENTL    Define extent for the read
FBACW2   CCW0   LOC_DATA,LOCATE,CC,LOCATEL    Establish location for read
FBACW3   CCW0   READ,0,0,0                    Read the directed record
*        CCW0   NOP,0,0,1                      ..then a NOP.
* If the preceding NOP CCW command is enabled, then the FBACW3 must set
* command chaining in the flag byte, setting the third operand to X'40'
        SPACE 1
* FBA extent used for reading a directed record. Unlike typical operations
* where the extent is constant and the locate data changes, when reading
* directed boot loader records, the extent changes and the locate information
* remains unchanged.
EXTENT   DC     XL4'40000000'    Extent file mask: Inhibit all writes
FRSTPSEC DC     FL4'0'           ** Physical first sector of the extent
FRSTLSEC DC     FL4'0'           First logical sector of the extent, always 0
* Last logical sector of the extent, always the same based upon record length
ENDLSEC  DC     FL4'0'
EXTENTL  EQU     *-EXTENT        Length of an FBA extent (16 bytes)
* ** This field is adjusted for each read of a directed record.
        SPACE 1
* FBA locate used for reading a directed record
LOCATE   DC     XL1'06'          Read sector operation being performed
        DC     XL1'00'          Replication count ignored for read sector
LOCSECS  DC     HL2'0'          Number of sectors being read
        DC     FL4'0'          First sector (relative to the extent), always 0
LOCATEL  EQU     *-LOCATE        Length of the FBA locate information (8 bytes)

```

SATK User Guide

The entire read sequence of the directed record is identical to that used in the Hello World in program 3. Refer to program 3 for a detailed description of how the I/O itself works.

The only difference is the I/O device used for the I/O (the FBA DASD volume, X'110') and the channel program address (FBACCW1).

Loading the Directed Record

As with the I/O operations, loading of the directed record requires some basic initialization.

```
* Initialize the static portions of record loading
      LA      8,6(4)      Locate start of directed record's content
      LA      3,HWM      Directed records may not overwrite me!
      ...
HWM      EQU      *      Can not load any directed record lower than here
```

Registers are used to maintain the information for the loading of directed records. All directed records are read into storage at the same location: the boot loader input/output area. General register 4 already points to this area. Each FBA directed record contains a 6-byte header composed of two fields:

- the destination address of the record's content (the first four bytes), followed by
- the record's binary content length (the last two bytes).

The above LA instruction causes general register 8 to point to the start of the binary content. Because this does not change, the same value is used through the entire loading process.

The high-water mark (HWM) location is used to ensure the boot loader is not overwritten by any directed record's content. This location is an EQUATE at the end of the boot loader, marking its highest address.

Immediately following the successful reading of a directed record, it is checked to ensure the boot loader is safe.

```
* Move the directed record to its residence address
      L       10,0(4)      Destination address of record's content
      CLR     3,10         Will data from record overwrite boot loader?
      BH      OVRWRITE     ..Yes, HWM higher than load address, quit now!
      ...
OVRWRITE LPSW OVERWRIT Code 01C Overwriting boot loader
      ...
OVERWRIT DWAIT PGM=04,CMP=0,CODE=01C Trying to overwrite boot loader
```

Now that it is known that the directed record is safe to move, a MVCL is used to do that following register setup. This time the directed record's header length is used to initialize the MVCL source and destination lengths. (The lengths are the same because the whole content is moved, as opposed to only a portion of it with some padding.)

SATK User Guide

```

ICM    9,B'0011',4(4)    Size of record being loaded
LR     11,9              Same size to the receiving location
LR     6,9              Increment for cumulative program loaded
MVCL   10,8             Move directed record
BC     B'0001',DESTRT    Destructive overlap detected (CC=3), quit
...
DESTRT LPSW DESTOVLP Code 020 Destructive overlap detected by MVCL
...
DESTOVLP DWAIT PGM=04,CMP=0,CODE=020 Destructive overlap detected by MVCL

```

With the successful loading of the directed record's content, the boot loader can increase the size of the loader program for the content contributed by this record. The size of the loaded content was preserved in general register 6 just before moving the directed record's content.

```

AL     6,LOD1BPLD        Increment loaded program size
ST     6,LOD1BPLD        Update the cumulative program size in LOD1

```

Now that the directed record has been successfully loaded to memory, it can be determined whether there are any more directed records to be read. The high-order bit (bit 0) of the directed record's address indicates whether this is the last directed record or not. A 1 in bit 0 indicate that the last directed record has been processed.

```

TM     0(4),X'80'        Bit 0 of destination address one?
BO     CKSIZE            ..Yes, check if correct amount loaded

```

If the branch is not taken, then additional directed records remain to be processed.

```

LR     6,5              Fetch the number of sectors read
AL     6,FRSTPSEC        Update starting sector for the next extent
ST     6,FRSTPSEC EXT - Update the FBA extent with new starting sector
B      READLOOP          Read the next record.

```

The FBA extent is updated for the next directed record to be read and processing of directed records continues.

If on the other hand, the above branch was taken, then all of the booted program's directed records have been loaded to memory. The last thing to check is whether the size of the booted program that was loaded matches the size expected for the booted program as understood by `iplasma.py`.

```

CKSIZE CLC LOD1BPLN,LOD1BPLD Do the cumulative sizes match in LOD1
      BNE CUMERROR          ..No, something went wrong, quit
...
CUMERROR LPSW BADSIZE Code 024 Cumulative booted program sizes do no match
...
BADSIZE DWAIT PGM=04,CMP=0,CODE=024 Cumulative booted program sizes mismatch

```

Transferring Control to the Booted Program

Now that the booted program was successfully loaded, the boot loader can transfer control to the booted program. LOD1 informs the boot loader where the booted program should be entered and in

SATK User Guide

what addressing mode. The S/370 boot loader can not change the addressing mode, so if either mode is set in the booted program's entry address, an error results.

```
TM      LOD1ENTR,X'80'      Is 31-bit addressing set in address?
BO      AMERROR              ..Yes, can not do that, so quit
TM      LOD1ENTR,X'01'      Is 64-bit addressing set in address?
BO      AMERROR              ..Yes, can not do that either, quit
...
AMERROR LPSW NOAMCHNG Code 028 Can not change addressing mode for booted pgm
...
NOAMCHNG DWAIT PGM=04,CMP=0, CODE=028 Can not change booted program's AMODE
```

Now that it is known that the booted program can be safely entered, the boot loader does its last duty to the boot loading process. It enters the booted program where `iplasma.py` tells it to:

```
L      15,LOD1ENTR          Fetch entry point for booted program from LOD1
BR      15                  Enter the booted program
```

Step 2 – Create IPL Medium

Unlike PGM3, PGM4 uses two list-directed IPL directories when creating the IPL medium. The main program, this time the booted program, remains in the `ldipl` directory. As with PGM3, the `ldipl` directory contains the bare-metal “Hello World” program.

The additional directory, `boot`, contains the boot loader used to load the booted program. Additional information is provided to `iplasma.py` when creating the IPL medium.

```
${IPLASMA} -v -f ld -m pgm4.3310 --records --asa=ASAREGN.bin \
--volser=PGM4 --boot boot/boot4.txt --lasa=ASAREGN.bin --recl 512 \
ldipl/pgm4.txt 2>&1 |tee iplasma-${sfx}.txt
```

As is readily seen, the additional information is supplied to `iplasma.py` for the additional directory, shown above in **bold** font. The `--boot` argument tells `iplasma.py` that the main program is to be booted by the program with this list-directed control file. Additionally, because it does not use the default file name for the assigned storage region, that file name is supplied by the `--lasa` argument.

As a consequence of these arguments, two directories are described as input to `iplasma.py`.

```
IPL program - Boot Loader: boot/boot4.txt
PSW      None
ASA      000000 ASAREGN.bin   Length: 512 (0x200)
Load     000400 BOOT4.bin     Length: 480 (0x1E0)

Booted program - list-directed source: ldipl/pgm4.txt
PSW      000000 ASAPSW        Length: 8 (0x8)
ASA      000000 ASAREGN.bin   Length: 512 (0x200)  IGNORED
```

SATK User Guide

Load 002000 BOOTED4.bin Length: 272 (0x110)

The boot loader is treated as the target of the IPL process and the “Hello World” program as the booted program.

The FBA DASD map must also be adjusted.

```
FBA DASD Map:
IPL0: 0-0
VOLLBL: 1-1
ASA: 2-2
LOD1: 3-3
BOOT4.bin: 4-4
BOOTED: 5-5
```

Two entries have been added, the **BOOT4.bin** allocation, for the boot loader and the **LOD1** record used by the boot loader. The “Hello World” program is identified as the **BOOTED** program and has its own allocation on the FBA volume.

The contents of the **LOD1** record are also provided. This is a series of fields with their assigned storage area addresses. Most of the fields are zero in this case and many are supplied by the boot loader itself when it executes (which of course has not happened yet).

```
LOD1 IPL Record 4:
000240 D3D6C4F1 LOD1 Record ID
000244 12 IPL medium information
000245 00 Boot loader flags
000246 0200 Maximum length of boot directed records in bytes
000248 00000110 Cumulative length of booted program on medium in bytes
00024C 00000000 Boot Loader Supplied: cumulative length of loaded program in bytes
000250 00002000 Booted program's entry address
000254 00 Boot Loader Supplied: Boot Loader's operating environment
000255 00 Boot Loader Supplied: Boot Loader's I/O architecture and mode
000256 00 Boot Loader Supplied: Boot Loader services
000257 00 Boot Loader Supplied: Booted program entry addressing mode
000258 00000005 Booted program starting physical FBA sector number
00025C 0000 Booted program starting CKD cylinder number
00025E 0000 Booted program starting CKD track (head) number
000260 00 Booted program starting CKD record number
000261 00 Number of CKD directed records per track
000262 0000 Maximum CKD cylinder number
000264 0000 Maximum CKD track (head) number
000266 0000 Boot Loader Supplied: Device Number of IPL subchannel
000268 00000000 Boot Loader Supplied: I/O address of IPL device
00026C 00002370 Boot Loader I/O area starting address
000270 00000000 Boot Loader Supplied: Boot Loader services address
000274 00000000 RESERVED
000278 00000000 RESERVED
00027C 00000000 RESERVED
000280 00000000 RESERVED
000284 00000000 RESERVED
000288 00000000 RESERVED
00028C 00000000 RESERVED
```

A few of the fields deserve some comment here. X'12' is the identifier for an FBA DASD IPL device. This indicates the type of medium created for the boot loader's use. The boot loader should validate

SATK User Guide

that this device type is supported by the boot loader. While the Guide programs each are self contained, this design is used in the case that multiple boot loaders are available and the reused loader may not support the device built by `iplasma.py`. Stuff happens.

X'0200' specifies the maximum size of a directed record. In this case, 512 is a single physical sector. It is derived from the `--rec l` argument to `iplasma.py`. This value is used in the channel command that reads the directed record. Some devices will require the incorrect length indication to be suppressed in the channel command and some will not.

The cumulative length of the booted program is 272 bytes, X'110'. As the boot loader reads and loads each directed record it will accumulate its own size in the following four-byte field at address X'24C'. The boot loader will compare its size with that supplied by `iplasma.py`. If the two sizes are different, the boot loader will enter a disabled wait state, indicating a boot failure.

The first physical sector is communicated to the boot loader by the **LOD1** field at X'258'. The FBA boot loader will continue reading sectors until the last directed record is flagged by bit 0 of the directed record being set to 1.

The last **LOD1** field worth discussion at this point is the boot loader I/O area address, at X'26C'. This location is positioned higher in memory than the booted program resides. The boot loader will use this location in the channel command that reads directed records as its starting address.

Each new component also must reside in storage. As no surprise, they are also reflected in the memory map.

Memory Map:

```
ASA: 000000-0001FF
IPL0: 000000-000017
CCW0: 000008-000017
LOD1: 000200-0003FF
BOOT4.bin: 000400-0005DF
BOOTED: 002000-00210F
IPL1: 002328-00236F
```

In addition to the added locations used by the boot loader, the “Hello World” program has moved to address X'2000' as required when using a boot loader.

Sector 5 is the single directed record that is loaded for the booted program.

```
FBA sector 5
000000 80002000 011005C0 D2070000 C0B6B722
000010 C0864810 C08E9D00 10004740 C06E4720
...
```

The first six bytes of the directed record constitute the record's header. As revealed (and not a surprise), the record is loaded in storage starting at address X'2000' for X'110' (272) bytes. The first bit being set to 1 (the starting 8 in the sector dump) tells the boot loader that this is also the last record

SATK User Guide

to be loaded for the booted program. The boot loader will increment its cumulative count by the number of bytes loaded, 272, making the loaded count 272. The two values, the one from `iplasma.py` and the count by the boot loader match, suggesting the entire program and nothing extra was loaded.

Step 3 – Run the Booted Program

The `ipl` script will, as before, executes the program. This script has been adjusted for `pgm4` file names.

The Hercules configuration file contains these statements:

```
# Hercules sample configuration file for pgm4.asm
ARCHMODE S/370          # S/370 targeted by pgm4.asm
MAINSIZE 64K             # Sample starts at X'2000'
NUMCPU 1
# DIAG8CMD enable       # Allow use of DIAG X'008'

# Devices
000F 3215-C /           # Console required by mainframe systems
0110 3310 pgm4.3310     # FBA IPL disk
```

The comment on `MAINSIZE` has been adjusted to reflect the new location of the booted program, `X'2000'`. The file names have been changed for the `pgm4` guide directory.

The Hercules run control file, `pgm4.rc`, is as follows.

```
# This file is intended for use with Hercules Hyperion
# Comment any of the tracing statements if you do not want the trace
t+
t+00F
t+110
# Perform the FBA based boot loading and execution of the "Hello World" program
ipl 110
```

Both the console, `X'00F'`, and the FBA DASD, `X'110'` have been enabled for tracing.

And, happily, in the Hercules log file is...

```
06:34:26 /Hello Bare-Metal World!
```

and

```
06:34:26 HHC00809I Processor CP00: disabled wait state 000A0000 00000000
```

Success!! The booted program worked.